

```
In [4]: import os
import gdown
import glob
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm

import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
import torch.optim as optim
import torchvision.transforms as transforms

from skimage.metrics import structural_similarity as ssim_metric, peak_signal
```

```
In [ ]: import zipfile
file_id = "1uJmDZw649XS-r-dYs9WD-0PwF_TIroVw"
output = "dataset.zip"
gdown.download(f"https://drive.google.com/uc?id={file_id}", output, quiet =

with zipfile.ZipFile("dataset.zip", "r") as zip_ref:
    zip_ref.extractall("/content/dataset")
```

```
Downloading...
From (original): https://drive.google.com/uc?id=1uJmDZw649XS-r-dYs9WD-0PwF_T
IroVw
From (redirected): https://drive.google.com/uc?id=1uJmDZw649XS-r-dYs9WD-0PwF
_TIroVw&confirm=t&uuiid=22729b50-4fdf-4c1c-b3c3-082e85210db3
To: /content/dataset.zip
100%|██████████| 533M/533M [00:06<00:00, 81.6MB/s]
```

```
In [6]: class LensingDataset(Dataset):
    def __init__(self, root_dir, transform = None):
        super(LensingDataset, self).__init__()
        self.root_dir = root_dir
        self.hr_dir = os.path.join(root_dir, 'HR')
        self.lr_dir = os.path.join(root_dir, 'LR')
        self.lr_files = sorted(glob.glob(os.path.join(self.lr_dir, '*.npz')))
        self.hr_files = sorted(glob.glob(os.path.join(self.hr_dir, '*.npz')))
        self.transform = transform

    def __len__(self):
        return len(self.lr_files)

    def __getitem__(self, idx):
        lr_image = np.load(self.lr_files[idx])
        hr_image = np.load(self.hr_files[idx])

        lr_image = torch.from_numpy(lr_image).float()
        hr_image = torch.from_numpy(hr_image).float()

        if self.transform:
            lr_image = self.transform(lr_image)
            hr_image = self.transform(hr_image)
```

```

lr_image = lr_image.repeat(3, 1, 1)
hr_image = hr_image.repeat(3, 1, 1)
return lr_image, hr_image

```

```

In [10]: class ResidualBlock(nn.Module):
    def __init__(self, n_feats, kernel_size=3, res_scale=0.1):
        super(ResidualBlock, self).__init__()
        self.res_scale = res_scale
        self.block = nn.Sequential(
            nn.Conv2d(n_feats, n_feats, kernel_size, padding=kernel_size//2),
            nn.ReLU(inplace=True),
            nn.Conv2d(n_feats, n_feats, kernel_size, padding=kernel_size//2)
        )

    def forward(self, x):
        res = self.block(x)
        return x + res * self.res_scale

```

```

In [11]: class EDSR(nn.Module):
    def __init__(self, scale_factor=2, n_resblocks=16, n_feats=64, res_scale=0.1):
        super(EDSR, self).__init__()
        self.head = nn.Conv2d(n_colors, n_feats, kernel_size=3, padding=1)

        self.body = nn.Sequential(
            *[ResidualBlock(n_feats, kernel_size=3, res_scale=res_scale) for _ in range(n_resblocks)]
        )
        self.body_conv = nn.Conv2d(n_feats, n_feats, kernel_size=3, padding=1)

        self.tail = nn.Sequential(
            nn.Conv2d(n_feats, n_feats * (scale_factor ** 2), kernel_size=3),
            nn.PixelShuffle(scale_factor),
            nn.Conv2d(n_feats, n_colors, kernel_size=3, padding=1)
        )

    def forward(self, x):
        x = self.head(x)
        res = self.body(x)
        res = self.body_conv(res)
        x = x + res
        x = self.tail(x)
        return x

```

```

In [13]: scale = 2
print("Initializing EDSR model...")
model = EDSR(scale_factor=scale, n_resblocks=8, n_feats=64, res_scale=0.1, n_colors=3)
model

```

Initializing EDSR model...

```

Out[13]: EDSR(
  (head): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (body): Sequential(
    (0): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (1): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (2): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (3): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (4): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (5): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
  )
)

```

```

    )
    (6): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (7): ResidualBlock(
      (block): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (1): ReLU(inplace=True)
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      )
    )
    (body_conv): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=
(1, 1))
    (tail): Sequential(
      (0): Conv2d(64, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): PixelShuffle(upscale_factor=2)
      (2): Conv2d(64, 3, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    )
  )
)

```

```

In [14]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
         model = model.to(device)

```

```

In [15]: criterion = nn.MSELoss()
         optimizer = optim.Adam(model.parameters(), lr=1e-4)

```

```

In [16]: root_dir = "/content/dataset/Dataset"
         train_dataset = LensingDataset(root_dir)

```

```

In [17]: train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True, num_wor

```

```

In [18]: num_epochs = 10
         print("Starting training on {} for {} epochs...".format(device, num_epochs))
         for epoch in range(num_epochs):
             model.train()
             running_loss = 0.0
             for lr_imgs, hr_imgs in tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_e
                 lr_imgs = lr_imgs.to(device)
                 hr_imgs = hr_imgs.to(device)

                 optimizer.zero_grad()
                 outputs = model(lr_imgs)
                 loss = criterion(outputs, hr_imgs)
                 loss.backward()
                 optimizer.step()

```

```

        running_loss += loss.item() * lr_imgs.size(0)

    epoch_loss = running_loss / len(train_dataset)
    print("Epoch [{}/{}], Loss: {:.6f}".format(epoch+1, num_epochs, epoch_lo

```

Starting training on cuda for 10 epochs...

Epoch 1/10: 100%|██████████| 1250/1250 [01:18<00:00, 15.91it/s]

Epoch [1/10], Loss: 0.000168

Epoch 2/10: 100%|██████████| 1250/1250 [01:18<00:00, 15.86it/s]

Epoch [2/10], Loss: 0.000060

Epoch 3/10: 100%|██████████| 1250/1250 [01:20<00:00, 15.48it/s]

Epoch [3/10], Loss: 0.000060

Epoch 4/10: 100%|██████████| 1250/1250 [01:20<00:00, 15.49it/s]

Epoch [4/10], Loss: 0.000059

Epoch 5/10: 100%|██████████| 1250/1250 [01:20<00:00, 15.47it/s]

Epoch [5/10], Loss: 0.000059

Epoch 6/10: 100%|██████████| 1250/1250 [01:21<00:00, 15.40it/s]

Epoch [6/10], Loss: 0.000059

Epoch 7/10: 100%|██████████| 1250/1250 [01:20<00:00, 15.45it/s]

Epoch [7/10], Loss: 0.000059

Epoch 8/10: 100%|██████████| 1250/1250 [01:21<00:00, 15.42it/s]

Epoch [8/10], Loss: 0.000059

Epoch 9/10: 100%|██████████| 1250/1250 [01:21<00:00, 15.41it/s]

Epoch [9/10], Loss: 0.000059

Epoch 10/10: 100%|██████████| 1250/1250 [01:21<00:00, 15.43it/s]

Epoch [10/10], Loss: 0.000059

```

In [19]: def evaluate(model, dataloader, device):
    model.eval()
    mse_total, ssim_total, psnr_total = 0.0, 0.0, 0.0
    count = 0
    with torch.no_grad():
        for lr_imgs, hr_imgs in tqdm(dataloader, desc="Evaluating"):
            lr_imgs = lr_imgs.to(device)
            hr_imgs = hr_imgs.to(device)

            outputs = model(lr_imgs)
            outputs_np = outputs.cpu().numpy()
            hr_imgs_np = hr_imgs.cpu().numpy()

            for i in range(outputs_np.shape[0]):
                # Convert to grayscale by taking the first channel.
                sr = outputs_np[i, 0, :, :]
                hr = hr_imgs_np[i, 0, :, :]
                mse_val = np.mean((sr - hr) ** 2)
                ssim_val = ssim_metric(sr, hr, data_range=hr.max()-hr.min())
                psnr_val = psnr_metric(hr, sr, data_range=hr.max()-hr.min())

                mse_total += mse_val
                ssim_total += ssim_val
                psnr_total += psnr_val
                count += 1

```

```
return mse_total/count, ssim_total/count, psnr_total/count
```

```
In [20]: mse_val, ssim_val, psnr_val = evaluate(model, train_loader, device)
print("Evaluation Metrics:")
print("MSE: {:.6f}".format(mse_val))
print("SSIM: {:.6f}".format(ssim_val))
print("PSNR: {:.6f}".format(psnr_val))
```

```
Evaluating: 100%|██████████| 1250/1250 [00:48<00:00, 25.52it/s]
Evaluation Metrics:
MSE: 0.000059
SSIM: 0.977661
PSNR: 42.343273
```

```
In [21]: def show_generated_images(model, dataloader, device, num_images=5):
        """
        Displays a few super-resolved images alongside the original LR and HR in
        """
        model.eval()
        lr_imgs, hr_imgs = next(iter(dataloader))
        lr_imgs = lr_imgs.to(device)
        hr_imgs = hr_imgs.to(device)

        with torch.no_grad():
            sr_imgs = model(lr_imgs)

        lr_imgs = lr_imgs.cpu().numpy()
        sr_imgs = sr_imgs.cpu().numpy()
        hr_imgs = hr_imgs.cpu().numpy()

        num_images = min(num_images, lr_imgs.shape[0])
        fig, axes = plt.subplots(num_images, 3, figsize=(12, 4 * num_images))

        if num_images == 1:
            axes = np.expand_dims(axes, axis=0)

        for i in range(num_images):
            lr_img = lr_imgs[i, 0, :, :]
            sr_img = sr_imgs[i, 0, :, :]
            hr_img = hr_imgs[i, 0, :, :]

            axes[i, 0].imshow(lr_img, cmap='gray')
            axes[i, 0].set_title("Low Resolution")
            axes[i, 0].axis('off')


            axes[i, 1].imshow(sr_img, cmap='gray')
            axes[i, 1].set_title("Super Resolved")
            axes[i, 1].axis('off')

            axes[i, 2].imshow(hr_img, cmap='gray')
            axes[i, 2].set_title("High Resolution")
            axes[i, 2].axis('off')
```


```
plt.tight_layout()  
plt.show()
```

```
In [22]: show_generated_images(model, train_loader, device, num_images=5)
```


Low Resolution




Super Resolved




High Resolution




Low Resolution




Super Resolved




High Resolution




Low Resolution




Super Resolved




High Resolution




Low Resolution




Super Resolved




High Resolution




Low Resolution





Super Resolved



High Resolution





```
In [23]: torch.save(model.state_dict(), 'edsr_strong_lensing_finetuned.pth')
         print("Model saved.")
```

Model saved.

```
In [24]: file_id = "1pLYfM-jFJT7TbTMVssuCCFvLzGdxMQ4h"
         output = "datasetB.zip"
         gdown.download(f"https://drive.google.com/uc?id={file_id}", output, quiet =
         with zipfile.ZipFile("datasetB.zip", "r") as zip_ref:
             zip_ref.extractall("/content/datasetB")
```

Downloading...

From: <https://drive.google.com/uc?id=1pLYfM-jFJT7TbTMVssuCCFvLzGdxMQ4h>

To: /content/datasetB.zip

100%|██████████| 21.6M/21.6M [00:00<00:00, 67.1MB/s]

```
In [25]: path = "/content/datasetB/Dataset/HR/HR_1.npy"
         image = np.load(path)
         image.shape
```

Out[25]: (1, 128, 128)

```
In [26]: path = "/content/datasetB/Dataset/LR/LR_1.npy"
         image = np.load(path)
         image.shape
```

Out[26]: (1, 64, 64)

```
In [37]: class RealLensingDataset(Dataset):
         def __init__(self, root_dir, transform = None):
             super(RealLensingDataset, self).__init__()
             self.lr_dir = os.path.join(root_dir, 'LR')
             self.hr_dir = os.path.join(root_dir, 'HR')
             self.lr_files = sorted(glob.glob(os.path.join(self.lr_dir, '*.npy')))
             self.hr_files = sorted(glob.glob(os.path.join(self.hr_dir, '*.npy')))
             self.transform = transform

         def __len__(self):
             return len(self.lr_files)

         def __getitem__(self, idx):
             lr_img = np.load(self.lr_files[idx])
             hr_img = np.load(self.hr_files[idx])

             lr_img = torch.from_numpy(lr_img).float()
             hr_img = torch.from_numpy(hr_img).float()

             lr_img = lr_img.repeat(3, 1, 1)
             hr_img = hr_img.repeat(3, 1, 1)
```

```

    if self.transform:
        lr_img = self.transform(lr_img)
        hr_img = self.transform(hr_img)

    return lr_img, hr_img

```

```

In [38]: scale = 2
        print("Initializing EDSR model for fine-tuning...")
        model = EDSR(scale_factor=scale, n_resblocks=8, n_feats=64, res_scale=0.1, r
Initializing EDSR model for fine-tuning...

```

```

In [39]: pretrained_weights = 'edsr_strong_lensing_finetuned.pth'
        model.load_state_dict(torch.load(pretrained_weights, map_location='cpu'))
        print("Pre-trained weights loaded.")
Pre-trained weights loaded.

```

<ipython-input-39-95e6a2d481d0>:2: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```

        model.load_state_dict(torch.load(pretrained_weights, map_location='cpu'))

```

```

In [40]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
        model = model.to(device)

```

```

In [41]: criterion = nn.MSELoss()
        optimizer = optim.Adam(model.parameters(), lr=1e-5)

```

```

In [45]: import random
        class RandomFlipRotate(object):
            """Apply random horizontal/vertical flips and random 90° rotations."""
            def __init__(self, p=0.5):
                self.p = p

            def __call__(self, image):
                if random.random() < self.p:
                    image = torch.flip(image, dims=[2]) # horizontal flip
                if random.random() < self.p:
                    image = torch.flip(image, dims=[1]) # vertical flip
                rotations = random.randint(0, 3)
                image = torch.rot90(image, k=rotations, dims=[1, 2])
                return image

```

```

In [43]: root_dir = "/content/datasetB/Dataset"
        data_transform = transforms.Compose([RandomFlipRotate(p=0.5)])

```



```

Epoch [1/10], Fine-tuning Loss: 0.005273
Fine-tune Epoch 2/10: 100%|██████████| 75/75 [00:01<00:00, 41.05it/s]
Epoch [2/10], Fine-tuning Loss: 0.002429
Fine-tune Epoch 3/10: 100%|██████████| 75/75 [00:02<00:00, 35.98it/s]
Epoch [3/10], Fine-tuning Loss: 0.002132
Fine-tune Epoch 4/10: 100%|██████████| 75/75 [00:01<00:00, 39.50it/s]
Epoch [4/10], Fine-tuning Loss: 0.002031
Fine-tune Epoch 5/10: 100%|██████████| 75/75 [00:01<00:00, 41.31it/s]
Epoch [5/10], Fine-tuning Loss: 0.002033
Fine-tune Epoch 6/10: 100%|██████████| 75/75 [00:02<00:00, 36.07it/s]
Epoch [6/10], Fine-tuning Loss: 0.002003
Fine-tune Epoch 7/10: 100%|██████████| 75/75 [00:01<00:00, 38.54it/s]
Epoch [7/10], Fine-tuning Loss: 0.001978
Fine-tune Epoch 8/10: 100%|██████████| 75/75 [00:01<00:00, 40.85it/s]
Epoch [8/10], Fine-tuning Loss: 0.001985
Fine-tune Epoch 9/10: 100%|██████████| 75/75 [00:01<00:00, 41.25it/s]
Epoch [9/10], Fine-tuning Loss: 0.001930
Fine-tune Epoch 10/10: 100%|██████████| 75/75 [00:01<00:00, 40.60it/s]
Epoch [10/10], Fine-tuning Loss: 0.001934

```

```

In [47]: def evaluate(model, dataloader, device):
          model.eval()
          mse_total, ssim_total, psnr_total = 0.0, 0.0, 0.0
          count = 0
          with torch.no_grad():
              for lr_imgs, hr_imgs in tqdm(dataloader, desc="Evaluating"):
                  lr_imgs = lr_imgs.to(device)
                  hr_imgs = hr_imgs.to(device)
                  outputs = model(lr_imgs)
                  outputs_np = outputs.cpu().numpy()
                  hr_imgs_np = hr_imgs.cpu().numpy()

                  for i in range(outputs_np.shape[0]):
                      # Convert to grayscale by taking the first channel.
                      sr = outputs_np[i, 0, :, :]
                      hr = hr_imgs_np[i, 0, :, :]
                      mse_val = np.mean((sr - hr)**2)
                      ssim_val = ssim_metric(sr, hr, data_range=hr.max()-hr.min())
                      psnr_val = psnr_metric(hr, sr, data_range=hr.max()-hr.min())

                      mse_total += mse_val
                      ssim_total += ssim_val
                      psnr_total += psnr_val
                      count += 1
          return mse_total/count, ssim_total/count, psnr_total/count

```

```

In [48]: mse_val, ssim_val, psnr_val = evaluate(model, real_loader, device)
          print("Evaluation Metrics on Real Data:")
          print("MSE: {:.6f}".format(mse_val))
          print("SSIM: {:.6f}".format(ssim_val))
          print("PSNR: {:.6f}".format(psnr_val))

```

```

Evaluating: 100%|██████████| 75/75 [00:01<00:00, 55.58it/s]

```

Evaluation Metrics on Real Data:

MSE: 0.001903

SSIM: 0.757113

PSNR: 29.645286

```
In [49]: def show_generated_images(model, dataloader, device, num_images=5):
        """
        Displays LR input, SR output, and HR ground truth images side by side.
        """
        model.eval()
        lr_imgs, hr_imgs = next(iter(dataloader))
        lr_imgs = lr_imgs.to(device)
        hr_imgs = hr_imgs.to(device)

        with torch.no_grad():
            sr_imgs = model(lr_imgs)

        lr_imgs = lr_imgs.cpu().numpy()
        sr_imgs = sr_imgs.cpu().numpy()
        hr_imgs = hr_imgs.cpu().numpy()

        num_images = min(num_images, lr_imgs.shape[0])
        fig, axes = plt.subplots(num_images, 3, figsize=(12, 4 * num_images))
        if num_images == 1:
            axes = np.expand_dims(axes, axis=0)

        for i in range(num_images):
            lr_img = lr_imgs[i, 0, :, :]
            sr_img = sr_imgs[i, 0, :, :]
            hr_img = hr_imgs[i, 0, :, :]

            axes[i, 0].imshow(lr_img, cmap='gray')
            axes[i, 0].set_title("Low Resolution")
            axes[i, 0].axis('off')

            axes[i, 1].imshow(sr_img, cmap='gray')
            axes[i, 1].set_title("Super Resolved")
            axes[i, 1].axis('off')

            axes[i, 2].imshow(hr_img, cmap='gray')
            axes[i, 2].set_title("High Resolution")
            axes[i, 2].axis('off')

        plt.tight_layout()
        plt.show()
```

```
In [50]: show_generated_images(model, real_loader, device, num_images=5)
```

