

```
In [1]: !pip install gdown
import gdown
import zipfile
file_id = "1znqUeFzYz-DeAE3dYXD17qoMPK82Whji"
output = "dataset.zip"
gdown.download(f"https://drive.google.com/uc?id={file_id}", output, quiet =

with zipfile.ZipFile("dataset.zip", "r") as zip_ref:
    zip_ref.extractall("/content/dataset")
```

Requirement already satisfied: gdown in /usr/local/lib/python3.11/dist-packages (5.2.0)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.11/dist-packages (from gdown) (4.13.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from gdown) (3.17.0)
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.11/dist-packages (from gdown) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from gdown) (4.67.1)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/dist-packages (from beautifulsoup4->gdown) (2.6)
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/python3.11/dist-packages (from beautifulsoup4->gdown) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests[socks]->gdown) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests[socks]->gdown) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests[socks]->gdown) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests[socks]->gdown) (2025.1.31)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/python3.11/dist-packages (from requests[socks]->gdown) (1.7.1)

Downloading...
From (original): https://drive.google.com/uc?id=1znqUeFzYz-DeAE3dYXD17qoMPK82Whji
From (redirected): https://drive.google.com/uc?id=1znqUeFzYz-DeAE3dYXD17qoMPK82Whji&confirm=t&uuid=920e2e13-edb0-44d5-8d5f-1f2d7f66783c
To: /content/dataset.zip
100%|██████████| 2.89G/2.89G [00:58<00:00, 49.3MB/s]

```
In [2]: file_id = "1uJmDZw649XS-r-dYs9WD-0PwF_TIroVw"
output = "datasetResolution.zip"

gdown.download(f"https://drive.google.com/uc?id={file_id}", output, quiet=False)

with zipfile.ZipFile(output, "r") as zip_ref:
    zip_ref.extractall("/content/datasetResolution")
```

```
Downloading...
From (original): https://drive.google.com/uc?id=1uJmDZw649XS-r-dYs9WD-0PwF_T
IroVw
From (redirected): https://drive.google.com/uc?id=1uJmDZw649XS-r-dYs9WD-0PwF
_TIroVw&confirm=t&uuid=c7024567-ef35-4a64-ad74-8882499d1db3
To: /content/datasetResolution.zip
100%|██████████| 533M/533M [00:12<00:00, 41.6MB/s]
```

```
In [3]: import os
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, random_split
from torchvision import transforms
from torch.cuda.amp import GradScaler, autocast
import torch.nn.functional as F
from tqdm import tqdm
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc, roc_auc_score
from skimage.metrics import structural_similarity as ssim
import math
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
from tqdm import tqdm
```

```
In [4]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

```
In [5]: class MAEDataset(Dataset):
    def __init__(self, file_paths, mask_ratio=0.85, patch_size=8, transform=
        self.file_paths = file_paths
        self.mask_ratio = mask_ratio
        self.patch_size = patch_size
        self.transform = transform

    def __len__(self):
        return len(self.file_paths)

    def __getitem__(self, idx):
        img = np.load(self.file_paths[idx])
        img = torch.tensor(img, dtype=torch.float32)
        # Adding channel dimension
        if img.ndim == 2:
            img = img.unsqueeze(0)
        if self.transform:
            img = self.transform(img)
        # Creating a masked version of the image.
        masked_img, mask = self.apply_mask(img)
        return masked_img, img, mask

    def apply_mask(self, img):
        """
```

```

"""
C, H, W = img.shape
patch_size = self.patch_size
n_patches_h = H // patch_size
n_patches_w = W // patch_size
total_patches = n_patches_h * n_patches_w
num_mask = int(total_patches * self.mask_ratio)
# Randomly choosing patch indices to mask
mask_indices = np.random.choice(total_patches, num_mask, replace=False)
mask = torch.zeros((n_patches_h, n_patches_w))
mask.view(-1)[mask_indices] = 1.0 # 1 indicates the patch is masked

# Creating a copy of the image and zero-out masked patches.
masked_img = img.clone()
for i in range(n_patches_h):
    for j in range(n_patches_w):
        if mask[i, j] == 1:
            masked_img[:, i*patch_size:(i+1)*patch_size, j*patch_size:(j+1)*patch_size] = 0
return masked_img, mask

```

```

In [6]: class ClassificationDataset(Dataset):
def __init__(self, data_path, classes, transform = None):
    self.samples = []
    for i, cls in enumerate(classes):
        for path in data_path[cls]:
            self.samples.append((path, i))
    self.transform = transform

def __len__(self):
    return len(self.samples)

def __getitem__(self, idx):
    path, label = self.samples[idx]
    img = np.load(path)
    img = torch.tensor(img, dtype=torch.float32)
    if img.ndim == 2:
        img = img.unsqueeze(0)
    if self.transform:
        img = self.transform(img)
    return img, label

```

```

In [7]: class SRDataset(Dataset):
def __init__(self, lr_paths, hr_paths, transform = None):
    self.lr_paths = lr_paths
    self.hr_paths = hr_paths
    self.transform = transform

def __len__(self):
    return len(self.lr_paths)

def __getitem__(self, idx):
    lr = np.load(self.lr_paths[idx])
    hr = np.load(self.hr_paths[idx])
    lr = torch.tensor(lr, dtype=torch.float32)
    hr = torch.tensor(hr, dtype = torch.float32)

```

```

if lr.ndim == 2:
    lr = lr.unsqueeze(0)
if hr.ndim == 2:
    hr = hr.unsqueeze(0)
if self.transform:
    lr = self.transform(lr)
    hr = self.transform(hr)
return lr, hr

```

```

In [8]: class SEBlock(nn.Module):
def __init__(self, channel, reduction = 16):
    super(SEBlock, self).__init__()
    self.avg_pool = nn.AdaptiveAvgPool2d(1)
    self.fc = nn.Sequential(
        nn.Linear(channel, channel // reduction, bias = False),
        nn.ReLU(inplace = True),
        nn.Linear(channel // reduction, channel, bias = False),
        nn.Sigmoid()
    )

def forward(self, x):
    b, c, _, _ = x.size()
    y = self.avg_pool(x).view(b, c)
    y = self.fc(y).view(b, c, 1, 1)
    return x * y

```

```

In [9]: class ResidualBlock(nn.Module):
def __init__(self, in_channels, out_channels, stride = 1, use_se = False):
    super(ResidualBlock, self).__init__()
    self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size = 3, stride = stride)
    self.bn1 = nn.BatchNorm2d(out_channels)
    self.act1 = nn.LeakyReLU(inplace = True)
    self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size = 3, stride = stride)
    self.bn2 = nn.BatchNorm2d(out_channels)
    self.act2 = nn.LeakyReLU(inplace = True)
    self.se = SEBlock(out_channels) if use_se else None

    self.downsample = None
    if stride != 1 or in_channels != out_channels:
        self.downsample = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size = 1, stride = stride),
            nn.BatchNorm2d(out_channels)
        )
    self.act2 = nn.LeakyReLU(inplace = True)

def forward(self, x):
    identity = x
    out = self.conv1(x)
    out = self.bn1(out)
    out = self.act1(out)
    out = self.conv2(out)
    out = self.bn2(out)
    if self.se is not None:
        out = self.se(out)
    if self.downsample is not None:

```

```

        identity = self.downsample(x)
        out += identity
        out = self.act2(out)
        return out

```

```

In [10]: class MAE(nn.Module):
    def __init__(self):
        super(MAE, self).__init__()
        self.encoder = nn.Sequential(
            ResidualBlock(1, 32, stride = 2, use_se = True),
            ResidualBlock(32, 64, stride = 2, use_se = True),
            ResidualBlock(64, 128, stride = 2, use_se = True),
            ResidualBlock(128, 256, stride = 2, use_se = True)
        )

        self.dc1_up = nn.ConvTranspose2d(256, 128, kernel_size = 3, stride = 2,
        self.dc1 = ResidualBlock(128 + 128, 128, stride = 1, use_se = True)
        self.dc2_up = nn.ConvTranspose2d(128, 64, kernel_size = 3, stride = 2,
        self.dc2 = ResidualBlock(64 + 64, 64, stride=1, use_se=True)
        self.dc3_up = nn.ConvTranspose2d(64, 32, kernel_size = 3, stride = 2,
        self.dc3 = ResidualBlock(32 + 32, 32, stride = 1, use_se = True)
        self.dc4_up = nn.ConvTranspose2d(32, 16, kernel_size = 3, stride = 2,
        self.dc4_conv = nn.Sequential(
            nn.Conv2d(16, 1, kernel_size = 3, stride = 1, padding = 1),
            nn.Sigmoid()
        )

    def forward(self, x):
        e1 = self.encoder[0](x)
        e2 = self.encoder[1](e1)
        e3 = self.encoder[2](e2)
        e4 = self.encoder[3](e3)

        d1_up = self.dc1_up(e4)
        d1_concat = torch.cat([d1_up, e3], dim = 1)
        d1 = self.dc1(d1_concat)

        d2_up = self.dc2_up(d1)
        d2_concat = torch.cat([d2_up, e2], dim = 1)
        d2 = self.dc2(d2_concat)

        d3_up = self.dc3_up(d2)
        d3_concat = torch.cat([d3_up, e1], dim = 1)
        d3 = self.dc3(d3_concat)

        d4_up = self.dc4_up(d3)
        out = self.dc4_conv(d4_up)
        return out, e4

```

```

In [11]: class Classifier(nn.Module):
    def __init__(self, encoder, num_classes = 2, latent_dim = 256):
        super(Classifier, self).__init__()
        self.encoder = encoder
        self.pool = nn.AdaptiveAvgPool2d((1, 1))

```

```

self.fc = nn.Linear(latent_dim, 2)

def forward(self, x):
    latent = self.encoder(x)
    x = self.pool(latent)
    x = torch.flatten(x, 1)
    logits = self.fc(x)
    return logits

```

```

In [12]: class CustomEncoder(nn.Module):
    def __init__(self, pretrained_mae):
        super(CustomEncoder, self).__init__()
        self.encoder = nn.Sequential(
            pretrained_mae.encoder[0],
            pretrained_mae.encoder[1]
        )
    def forward(self, x):
        return self.encoder(x)

```

```

In [ ]: class SRDecoderModified(nn.Module):
    def __init__(self, in_channels=64, target_size=(150, 150)):
        super(SRDecoderModified, self).__init__()
        self.target_size = target_size
        self.up1 = nn.ConvTranspose2d(in_channels, 32, kernel_size=4, stride=2, padding=1)
        self.conv1 = nn.Conv2d(32, 32, kernel_size=3, padding=1)
        self.act = nn.LeakyReLU(inplace=True)
        self.up2 = nn.ConvTranspose2d(32, 16, kernel_size=4, stride=2, padding=1)
        self.conv2 = nn.Conv2d(16, 16, kernel_size=3, padding=1)
        self.up3 = nn.ConvTranspose2d(16, 8, kernel_size=4, stride=2, padding=1)
        self.conv3 = nn.Conv2d(8, 1, kernel_size=3, padding=1)
        self.final_act = nn.Sigmoid()

    def forward(self, x):
        x = self.up1(x)
        x = self.conv1(x)
        x = self.act(x)
        x = self.up2(x)
        x = self.conv2(x)
        x = self.act(x)
        x = self.up3(x)
        x = self.conv3(x)
        if x.shape[2:] != self.target_size:
            diff_h = x.shape[2] - self.target_size[0]
            diff_w = x.shape[3] - self.target_size[1]
            x = x[:, :, diff_h//2:diff_h//2+self.target_size[0], diff_w//2:diff_w//2+self.target_size[1]]
        x = self.final_act(x)
        return x

```

```

In [14]: class SRModel(nn.Module):
    def __init__(self, pretrained_encoder, target_size=(150, 150)):
        super(SRModel, self).__init__()
        self.encoder = pretrained_encoder
        self.decoder = SRDecoderModified(in_channels=64, target_size=target_size)

    def forward(self, x):

```

```

        latent = self.encoder(x)
        out = self.decoder(latent)
        return out

```

```

In [15]: def train_mae(model, dataloader, optimizer, scheduler, device, scaler, epochs):
    model.train()
    criterion = nn.MSELoss()
    for epoch in range(epochs):
        epoch_loss = 0.0
        pbar = tqdm(dataloader, desc = f"[MAE Pre-training] Epoch {epoch+1}/{epochs}")
        for masked_img, img, _ in pbar:
            masked_img = masked_img.to(device)
            img = img.to(device)
            optimizer.zero_grad()
            with autocast():
                recon, _ = model(masked_img)
                loss = criterion(recon, img)
            scaler.scale(loss).backward()
            scaler.step(optimizer)
            scaler.update()
            epoch_loss += loss.item() * masked_img.size(0)
            pbar.set_postfix(loss=f"{loss.item():.4f}")
        scheduler.step()
        print(f"[MAE Pre-training] Epoch [{epoch+1}/{epochs}] Loss: {epoch_loss/epochs}")

```

```

In [16]: def train_classifier(model, dataloader, optimizer, scheduler, device, scaler, epochs):
    model.train()
    criterion = nn.CrossEntropyLoss()
    for epoch in range(epochs):
        epoch_loss = 0.0
        pbar = tqdm(dataloader, desc=f"[Classifier Fine-tuning] Epoch {epoch+1}/{epochs}")
        for imgs, labels in pbar:
            imgs = imgs.to(device)
            labels = labels.to(device)
            optimizer.zero_grad()
            with autocast():
                logits = model(imgs)
                loss = criterion(logits, labels)
            scaler.scale(loss).backward()
            scaler.step(optimizer)
            scaler.update()
            epoch_loss += loss.item() * imgs.size(0)
            pbar.set_postfix(loss=f"{loss.item():.4f}")
        scheduler.step()
        print(f"[Classifier Fine-tuning] Epoch [{epoch+1}/{epochs}] Loss: {epoch_loss/epochs}")

```

```

In [17]: def train_sr(model, dataloader, optimizer, device, scaler, epochs=10):
    model.train()
    criterion = nn.MSELoss()
    for epoch in range(epochs):
        epoch_loss = 0.0
        pbar = tqdm(dataloader, desc=f"SR Training Epoch {epoch+1}/{epochs}")
        for lr, hr in pbar:
            lr = lr.to(device)
            hr = hr.to(device)

```

```

optimizer.zero_grad()
with torch.amp.autocast("cuda"):
    sr = model(lr)
    loss = criterion(sr, hr)
    scaler.scale(loss).backward()
    scaler.step(optimizer)
    scaler.update()
    epoch_loss += loss.item() * lr.size(0)
    pbar.set_postfix(loss=f"{loss.item():.4f}")
print(f"Epoch {epoch+1}/{epochs}, SR Loss: {epoch_loss/len(dataloader)}")

```

```

In [18]: def evaluate_and_plot_roc(model, dataloader, device, num_classes=2):
    model.eval()
    all_labels = []
    all_probs = []

    with torch.no_grad():
        for imgs, labels in tqdm(dataloader, desc="Evaluating", leave=False):
            imgs = imgs.to(device)
            logits = model(imgs)
            if num_classes > 1:
                probs = torch.softmax(logits, dim=1)
            else:
                probs = torch.sigmoid(logits)
            all_probs.append(probs.cpu().numpy())
            all_labels.append(labels.cpu().numpy())

    # Concatenate lists into single arrays.
    all_probs = np.concatenate(all_probs, axis=0)
    all_labels = np.concatenate(all_labels, axis=0)

    # Ensure all_probs is 2D. If 1D, reshape to [n_samples, 1].
    if all_probs.ndim == 1:
        all_probs = all_probs[:, np.newaxis]
    if num_classes == 2 and all_probs.shape[1] == 1:
        all_probs = np.hstack([1 - all_probs, all_probs])
    if num_classes == 2:
        y_true = np.zeros((all_labels.shape[0], num_classes))
        y_true[np.arange(all_labels.shape[0]), all_labels] = 1
    else:
        y_true = label_binarize(all_labels, classes=list(range(num_classes)))

    # Compute ROC curve and AUC for each class.
    fpr = {}
    tpr = {}
    roc_auc = {}
    for i in range(num_classes):
        fpr[i], tpr[i], _ = roc_curve(y_true[:, i], all_probs[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])

    # Plot ROC curves for all classes.
    plt.figure(figsize=(8, 6))
    for i in range(num_classes):
        plt.plot(fpr[i], tpr[i], lw=2, label=f"Class {i} (AUC = {roc_auc[i]:.2f})")
    plt.plot([0, 1], [0, 1], 'k--', lw=2, label="False Positive Rate")

```



```

plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend(loc="lower right")
plt.show()

return all_labels, all_probs, roc_auc

```

```

In [19]: def evaluate_sr(model, dataloader, device):
    model.eval()
    mse_loss = 0.0
    psnr_total = 0.0
    ssim_total = 0.0
    criterion = nn.MSELoss(reduction='mean')
    n = 0
    with torch.no_grad():
        for lr, hr in tqdm(dataloader, desc="Evaluating SR", leave=False):
            lr = lr.to(device)
            hr = hr.to(device)
            sr = model(lr)
            mse = criterion(sr, hr).item()
            mse_loss += mse * lr.size(0)
            # Compute PSNR (assuming pixel values in [0,1])
            psnr_val = 10 * math.log10(1.0 / mse) if mse > 0 else 100
            psnr_total += psnr_val * lr.size(0)
            # Compute SSIM for each image (convert tensors to numpy arrays)
            sr_np = sr.cpu().numpy()
            hr_np = hr.cpu().numpy()
            for i in range(sr_np.shape[0]):
                # Assuming single channel images
                ssim_val = ssim(hr_np[i, 0], sr_np[i, 0], data_range=hr_np[i, 0].max() - hr_np[i, 0].min())
                ssim_total += ssim_val
                n += 1
    mse_avg = mse_loss / len(dataloader.dataset)
    psnr_avg = psnr_total / len(dataloader.dataset)
    ssim_avg = ssim_total / n
    print(f"SR Evaluation - MSE: {mse_avg:.4f}, PSNR: {psnr_avg:.2f}, SSIM:

```

```

In [20]: base_path_for_6A = "/content/dataset/Dataset"
classes = ['no_sub', 'cdm']
data_paths = {
    cls: [os.path.join(base_path_for_6A, cls, f)
          for f in os.listdir(os.path.join(base_path_for_6A, cls)) if f.endswith('.')]
    for cls in classes
}

```

```

In [21]: batch_size = 32
pretrain_epochs = 10
finetune_epochs = 10
lr = 1e-3
transform = None

```

```

In [22]: no_sub_paths = data_paths['no_sub']
pretrain_dataset = MAEDataset(no_sub_paths, mask_ratio=0.75, patch_size=8, transform=transform)
pretrain_loader = DataLoader(pretrain_dataset, batch_size=batch_size, shuffle=True)

```

```
/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624:
UserWarning: This DataLoader will create 4 worker processes in total. Our suggested
max number of worker in current system is 2, which is smaller than what this
DataLoader is going to create. Please be aware that excessive worker creation
might get DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
warnings.warn(
```

```
In [23]: modelMAE = MAE().to(device)
optimizer_mae = optim.Adam(modelMAE.parameters(), lr=lr)
scheduler_mae = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer_mae, T
scaler = GradScaler()
```

```
<ipython-input-23-6e79cb26ed3e>:4: FutureWarning: `torch.cuda.amp.GradScaler
(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)`
instead.
scaler = GradScaler()
```

```
In [24]: print("Starting MAE pre-training...")
train_mae(modelMAE, pretrain_loader, optimizer_mae, scheduler_mae, device, s
```

Starting MAE pre-training...

```
[MAE Pre-training] Epoch 1/10: 0%|          | 0/921 [00:00<?, ?it/s]<ipyth
on-input-15-e6c8494355bf>:11: FutureWarning: `torch.cuda.amp.autocast(arg
s...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` inste
ad.
```

```
    with autocast():
```

```
[MAE Pre-training] Epoch [1/10] Loss: 0.0619
```

```
[MAE Pre-training] Epoch [2/10] Loss: 0.0577
```

```
[MAE Pre-training] Epoch [3/10] Loss: 0.0576
```

```
[MAE Pre-training] Epoch [4/10] Loss: 0.0575
```

```
[MAE Pre-training] Epoch [5/10] Loss: 0.0575
```

```
[MAE Pre-training] Epoch [6/10] Loss: 0.0574
```

```
[MAE Pre-training] Epoch [7/10] Loss: 0.0574
```

```
[MAE Pre-training] Epoch [8/10] Loss: 0.0574
```

```
[MAE Pre-training] Epoch [9/10] Loss: 0.0574
```

```
[MAE Pre-training] Epoch [10/10] Loss: 0.0574
```

```
In [25]: classifier_model = Classifier(encoder=modelMAE.encoder, num_classes=len(class
optimizer_cls = optim.Adam(classifier_model.parameters(), lr=lr)
scheduler_cls = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer_cls, T
```

```
In [26]: class_dataset = ClassificationDataset(data_paths, classes, transform=transfo
train_size = int(0.8 * len(class_dataset))
val_size = len(class_dataset) - train_size
train_dataset, val_dataset = random_split(class_dataset, [train_size, val_si
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, r
```

```
In [27]: train_classifier(classifier_model, train_loader, optimizer_cls, scheduler_cl
```

```
[Classifier Fine-tuning] Epoch 1/10: 0%|          | 0/1481 [00:00<?, ?it/s]<ipython-input-16-d471aa5bcfld>:11: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
```

```
    with autocast():
```

```
[Classifier Fine-tuning] Epoch [1/10] Loss: 0.0986
```

```
[Classifier Fine-tuning] Epoch [2/10] Loss: 0.0096
```

```
[Classifier Fine-tuning] Epoch [3/10] Loss: 0.0063
```

```
[Classifier Fine-tuning] Epoch [4/10] Loss: 0.0043
```

```
[Classifier Fine-tuning] Epoch [5/10] Loss: 0.0027
```

```
[Classifier Fine-tuning] Epoch [6/10] Loss: 0.0021
```


```
[Classifier Fine-tuning] Epoch [7/10] Loss: 0.0013
```

```
[Classifier Fine-tuning] Epoch [8/10] Loss: 0.0006
```

```
[Classifier Fine-tuning] Epoch [9/10] Loss: 0.0001
```

```
[Classifier Fine-tuning] Epoch [10/10] Loss: 0.0000
```

```
In [28]: all_labels, all_probs, roc_auc = evaluate_and_plot_roc(classifier_model, val
```



```
In [29]: torch.save(modelMAE.state_dict(), 'pretrained_mae_weights.pth')
```

```
In [30]: base_path = "/content/datasetResolution/Dataset"
         ops.path.join(base_path, "LR")
```

```

hr_dir = os.path.join(base_path, "HR")
lr_paths = [os.path.join(lr_dir, f) for f in os.listdir(lr_dir) if f.endswith('l.png')]
hr_paths = [os.path.join(hr_dir, f) for f in os.listdir(hr_dir) if f.endswith('h.png')]
lr_paths.sort()
hr_paths.sort()

```

```
In [31]: sr_dataset = SRDataset(lr_paths, hr_paths, transform=None)
```

```
In [32]: train_size = int(0.8 * len(sr_dataset))
val_size = len(sr_dataset) - train_size
train_dataset, val_dataset = random_split(sr_dataset, [train_size, val_size])
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, num_workers=4)
val_loader = DataLoader(val_dataset, batch_size=16, shuffle=False, num_workers=4)
```

```

/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624:
UserWarning: This DataLoader will create 4 worker processes in total. Our suggested
max number of worker in current system is 2, which is smaller than what this
DataLoader is going to create. Please be aware that excessive worker creation
might get DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(

```

```
In [35]: pretrained_MAE = modelMAE
pretrained_MAE = pretrained_MAE.to(device)
pretrained_MAE.load_state_dict(torch.load('pretrained_mae_weights.pth'))
custom_encoder = CustomEncoder(pretrained_MAE)
sr_model = SRModel(pretrained_encoder=custom_encoder, target_size=(150,150))
```

```
In [36]: optimizer = optim.Adam(sr_model.parameters(), lr=1e-4)
scaler = torch.amp.GradScaler('cuda')
print("Starting Super-Resolution Fine-Tuning...")
train_sr(sr_model, train_loader, optimizer, device, scaler, epochs=10)
```

Starting Super-Resolution Fine-Tuning...

```

SR Training Epoch 1/10: 0%|          | 0/500 [00:00<?, ?it/s]/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning:
This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this
DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to
avoid potential slowness/freeze if necessary.
  warnings.warn(

```

Epoch 1/10, SR Loss: 0.0290

Epoch 2/10, SR Loss: 0.0005

Epoch 3/10, SR Loss: 0.0003

Epoch 4/10, SR Loss: 0.0003

Epoch 5/10, SR Loss: 0.0002

Epoch 6/10, SR Loss: 0.0002

Epoch 7/10, SR Loss: 0.0002

Epoch 8/10, SR Loss: 0.0002

Epoch 9/10, SR Loss: 0.0002

```
In [37]: evaluate_sr(sr_model, val_loader, device)
```

SR Evaluation - MSE: 0.0001, PSNR: 38.28, SSIM: 0.9439

```
In [40]: def compare_hr_sr(model, dataloader, device):
    model.eval()
    with torch.no_grad():
        for lr, hr in dataloader:
            lr = lr.to(device)
            hr = hr.to(device)
            sr = model(lr)

            # Iterate through each image in the batch
            for i in range(hr.shape[0]): # hr.shape[0] is the batch size
                hr_np = hr[i].cpu().numpy().squeeze(0) # Select image i from the
                sr_np = sr[i].cpu().numpy().squeeze(0) # Select image i from the

                ssim_score = ssim(hr_np, sr_np, data_range=hr_np.max()-hr_np.min())


                fig, axes = plt.subplots(1, 2, figsize = (10, 5))
                axes[0].imshow(hr_np, cmap='gray')
                axes[0].set_title('Original HR')
                axes[0].axis('off')
                axes[1].imshow(sr_np, cmap='gray')
                axes[1].set_title(f'Super-Resolved (SSIM: {ssim_score:.4f})')
                axes[1].axis('off')
                plt.show()
            break
```

```
In [41]: compare_hr_sr(sr_model, val_loader, device)
```


/usr/local/lib/python3.11/dist-packages/torch/utils/data/dataloader.py:624: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn(
```


Original HR




Super-Resolved (SSIM: 0.9477)




Original HR




Super-Resolved (SSIM: 0.9337)




Original HR




Super-Resolved (SSIM: 0.9429)




Original HR




Super-Resolved (SSIM: 0.9502)




Original HR




Super-Resolved (SSIM: 0.9446)




Original HR




Super-Resolved (SSIM: 0.9383)




Original HR




Super-Resolved (SSIM: 0.9441)




Original HR




Super-Resolved (SSIM: 0.9438)




Original HR




Super-Resolved (SSIM: 0.9455)




Original HR




Super-Resolved (SSIM: 0.9462)




Original HR




Super-Resolved (SSIM: 0.9437)




Original HR




Super-Resolved (SSIM: 0.9391)




Original HR




Super-Resolved (SSIM: 0.9434)




Original HR




Super-Resolved (SSIM: 0.9454)




Original HR




Super-Resolved (SSIM: 0.9418)



Original HR



Super-Resolved (SSIM: 0.9400)



This notebook was converted with convert.ploomber.io