

```
In [1]: !pip install gdown
```

```
Requirement already satisfied: gdown in /usr/local/lib/python3.11/dist-packa
ges (5.2.0)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.11/d
ist-packages (from gdown) (4.13.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-pa
ckages (from gdown) (3.17.0)
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.11/
dist-packages (from gdown) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packag
es (from gdown) (4.67.1)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.11/di
st-packages (from beautifulsoup4->gdown) (2.6)
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/py
thon3.11/dist-packages (from beautifulsoup4->gdown) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/py
thon3.11/dist-packages (from requests[socks]->gdown) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dis
t-packages (from requests[socks]->gdown) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.
11/dist-packages (from requests[socks]->gdown) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.
11/dist-packages (from requests[socks]->gdown) (2025.1.31)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/pyth
on3.11/dist-packages (from requests[socks]->gdown) (1.7.1)
```

```
In [2]: import gdown
import zipfile
file_id = "1ZEyNME043u3qhJAwJeBZxFBEYc_pVYZQ"
output = "dataset.zip"
gdown.download(f"https://drive.google.com/uc?id={file_id}", output, quiet =

with zipfile.ZipFile("dataset.zip", "r") as zip_ref:
    zip_ref.extractall("/content/dataset")
```

```
Downloading...
From (original): https://drive.google.com/uc?id=1ZEyNME043u3qhJAwJeBZxFBEYc_
pVYZQ
From (redirected): https://drive.google.com/uc?id=1ZEyNME043u3qhJAwJeBZxFBEY
c_pVYZQ&confirm=t&uuid=414d338e-b352-47d5-98f5-32e4038999bf
To: /content/dataset.zip
100%|██████████| 1.13G/1.13G [00:14<00:00, 79.2MB/s]
```

```
In [3]: import os
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
import torchvision.models as models
```

```
import torchvision.transforms as transforms
from tqdm import tqdm
```

```
In [4]: np.random.seed(42)
```

```
In [5]: class LensingDataset(Dataset):
    def __init__(self, data_paths, labels, transform = None):
        self.data_paths = data_paths
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.data_paths)

    def __getitem__(self, idx):
        img = np.load(self.data_paths[idx])
        if img.shape != (1, 150, 150):
            img = img.reshape(1, 150, 150)

        img_tensor = torch.from_numpy(img).float()

        if self.transform:
            img_tensor = self.transform(img_tensor)

        if img_tensor.shape[0] == 1:
            img_tensor = img_tensor.repeat(3, 1, 1)

        return img_tensor, self.labels[idx]
```

```
In [6]: def check_data_integrity(data_loader):
    try:
        inputs, labels = next(iter(data_loader))
        print(f"Data check - Input shape: {inputs.shape}, Labels shape: {labels.shape}")
        print(f"Data check - Input dtype: {inputs.dtype}, Labels dtype: {labels.dtype}")
        print(f"Data check - Labels unique values: {torch.unique(labels)}")
        return True
    except Exception as e:
        print(f"Data integrity check failed: {e}")
        return False
```

```
In [7]: def load_data(base_dir):
    class_folders = ['no', 'sphere', 'vort']
    data_paths = []
    labels = []

    for class_idx, folder in enumerate(class_folders):
        folder_path = os.path.join(base_dir, folder)

        if not os.path.exists(folder_path):
            print(f"Warning: Folder {folder_path} does not exist")
            continue

        npy_files = [f for f in os.listdir(folder_path) if f.endswith('.npy')]
        file_list = [os.path.join(folder_path, f) for f in npy_files]
        data_paths.extend(file_list)
```

```

labels.extend([class_idx] * len(file_list))

unique_labels, counts = np.unique(labels, return_counts=True)
print("Class distribution:")
for label, count in zip(unique_labels, counts):
    print(f"Class {label}: {count} samples")
return data_paths, labels

```

```

In [8]: def get_model(model_name, num_classes=3):
        if model_name == 'resnet50':
            model = models.resnet50(weights=models.ResNet50_Weights.IMAGENET1K_V1)
            model.fc = nn.Linear(model.fc.in_features, num_classes)
        elif model_name == 'efficientnet_b0':
            model = models.efficientnet_b0(weights=models.EfficientNet_B0_Weights.DEFAULT)
            model.classifier[1] = nn.Linear(model.classifier[1].in_features, num_classes)
        elif model_name == 'densenet121':
            model = models.densenet121(weights=models.DenseNet121_Weights.IMAGENET1K_V1)
            model.classifier = nn.Linear(model.classifier.in_features, num_classes)
        else:
            raise ValueError(f"Model {model_name} not supported")

        return model

```

```

In [18]: def train_model(model, train_loader, val_loader, criterion, optimizer, scheduler):
        best_model_wts = None
        best_auc = 0.0
        history = {
            'train_loss' : [], 'val_loss' : [],
            'train_acc' : [], 'val_acc' : []
        }
        if not check_data_integrity(train_loader) or not check_data_integrity(val_loader):
            print("Data integrity check failed. Aborting training.")
            return model, history

        for epoch in range(num_epochs):
            print(f'Epoch {epoch+1}/{num_epochs}')
            print('-' * 10)

            model.train()
            train_loss = 0.0
            train_corrects = 0

            for inputs, labels in tqdm(train_loader, desc = "Training"):
                try:
                    inputs = inputs.to(device)
                    labels = labels.to(device)

                    optimizer.zero_grad()

                    with torch.set_grad_enabled(True):
                        outputs = model(inputs)
                        _, preds = torch.max(outputs, 1)
                        loss = criterion(outputs, labels)

```

```

        loss.backward()
        optimizer.step()

    except Exception as e:
        print(f"Error during training: {e}")
        continue

    if scheduler:
        scheduler.step()

    train_loss += loss.item() * inputs.size(0)
    train_corrects += torch.sum(preds == labels.data)

    epoch_train_loss = train_loss / len(train_loader.dataset)
    epoch_train_acc = float(train_corrects) / len(train_loader.dataset)

    history['train_loss'].append(epoch_train_loss)
    history['train_acc'].append(epoch_train_acc)

    print(f'Train Loss: {epoch_train_loss:.4f} Acc: {epoch_train_acc:.4f}')

    model.eval()
    val_loss = 0.0
    val_corrects = 0
    all_labels = []
    all_probs = []

    for inputs, labels in tqdm(val_loader, desc="Validation"):
        try:
            inputs = inputs.to(device)
            labels = labels.to(device)

            with torch.no_grad():
                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)
                loss = criterion(outputs, labels)
                probs = nn.Softmax(dim = 1)(outputs)

            val_loss += loss.item() * inputs.size(0)
            val_corrects += torch.sum(preds == labels.data)

            all_labels.append(labels.cpu().numpy())
            all_probs.append(probs.cpu().numpy())

        except Exception as e:
            print(f"Error during validation: {e}")
            continue

    if len(all_labels) == 0 or len(all_probs) == 0:
        print("No valid validation batches. Skipping ROC calculation.")
        continue

    epoch_val_loss = val_loss / len(val_loader.dataset)
    epoch_val_acc = float(val_corrects) / len(val_loader.dataset)

```

```

history['val_loss'].append(epoch_val_loss)
history['val_acc'].append(epoch_val_acc)

print(f'Val Loss: {epoch_val_loss:.4f} Acc: {epoch_val_acc:.4f}')

try:
    all_labels = np.concatenate(all_labels)
    all_probs = np.concatenate(all_probs)

    unique_labels = np.unique(all_labels)

    mean_auc = 0
    class_count = 0

    for class_idx in range(3):
        if class_idx not in unique_labels and len(unique_labels) < 3:
            print(f"Warning: Class {class_idx} not in validation set")
            continue

        binary_labels = (all_labels == class_idx).astype(int)
        fpr, tpr, _ = roc_curve(binary_labels, all_probs[:, class_idx])
        roc_auc = auc(fpr, tpr)
        mean_auc += roc_auc
        class_count += 1

    mean_auc /= max(1, class_count)
    print(f'Mean AUC: {mean_auc:.4f}')

    if mean_auc > best_auc:
        best_auc = mean_auc
        best_model_wts = model.state_dict().copy()

except Exception as e:
    print(f"Error calculating ROC/AUC: {e}")
    continue

print(f'Best validation AUC: {best_auc:.4f}')

if best_model_wts is not None:
    model.load_state_dict(best_model_wts)
return model, history

```

```

In [19]: def evaluate_model(model, test_loader, device='cuda'):
    model.eval()
    all_labels = []
    all_probs = []
    all_preds = []

    with torch.no_grad():
        for inputs, labels in tqdm(test_loader, desc="Testing"):
            try:
                inputs = inputs.to(device)
                labels = labels.to(device)

                outputs = model(inputs)

```

```

_, preds = torch.max(outputs, 1)
probs = nn.Softmax(dim=1)(outputs)

all_labels.append(labels.cpu().numpy())
all_probs.append(probs.cpu().numpy())
all_preds.append(preds.cpu().numpy())

except Exception as e:
    print(f"Error during evaluation: {e}")
    continue

if not all_labels or not all_probs or not all_preds:
    print("No valid test batches. Cannot evaluate model.")
    return 0, 0, None

all_labels = np.concatenate(all_labels)
all_probs = np.concatenate(all_probs)
all_preds = np.concatenate(all_preds)

# Calculate accuracy
accuracy = np.mean(all_preds == all_labels)
print(f'Test Accuracy: {accuracy:.4f}')

# Plot ROC curves and calculate AUC
plt.figure(figsize=(10, 8))

class_names = ['No Substructure', 'Sphere Substructure', 'Vortex Substru
colors = ['blue', 'red', 'green']

mean_auc = 0
class_count = 0

# Get unique classes in test set
unique_labels = np.unique(all_labels)

for i, class_name in enumerate(class_names):
    # Skip classes that don't exist in the test set
    if i not in unique_labels and len(unique_labels) < 3:
        print(f"Warning: Class {i} not in test set")
        continue

    try:
        binary_labels = (all_labels == i).astype(int)
        fpr, tpr, _ = roc_curve(binary_labels, all_probs[:, i])
        roc_auc = auc(fpr, tpr)
        mean_auc += roc_auc
        class_count += 1

        plt.plot(fpr, tpr, color=colors[i], lw=2,
                 label=f'{class_name} (AUC = {roc_auc:.3f})')
    except Exception as e:
        print(f"Error calculating ROC for class {i}: {e}")
        continue

mean_auc /= max(1, class_count) # Avoid division by zero

```

```

# Plot diagonal line
plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curves (Mean AUC = {mean_auc:.3f})')
plt.legend(loc="lower right")
plt.grid(True)

return accuracy, mean_auc, plt

```

```

In [20]: def main(base_dir, model_name='efficientnet_b0', batch_size=16, num_epochs=100):
# Check if CUDA is available
if torch.cuda.is_available():
    device = torch.device("cuda")
    # Print CUDA information for debugging
    print(f"Using CUDA: {torch.cuda.get_device_name(0)}")
    print(f"CUDA device count: {torch.cuda.device_count()}")
    print(f"CUDA version: {torch.version.cuda}")
else:
    device = torch.device("cpu")
    print("CUDA not available, using CPU")

# Load data
try:
    data_paths, labels = load_data(base_dir)
except Exception as e:
    print(f"Error loading data: {e}")
    return None, 0, 0

# Split data
try:
    train_paths, test_paths, train_labels, test_labels = train_test_split(
        data_paths, labels, test_size=0.2, random_state=42, stratify=labels
    )

    train_paths, val_paths, train_labels, val_labels = train_test_split(
        train_paths, train_labels, test_size=0.2, random_state=42, stratify=train_labels
    )
except Exception as e:
    print(f"Error splitting data: {e}")
    return None, 0, 0

print(f"Train samples: {len(train_paths)}")
print(f"Validation samples: {len(val_paths)}")
print(f"Test samples: {len(test_paths)}")

# Create datasets and dataloaders
try:
    train_dataset = LensingDataset(train_paths, train_labels)
    val_dataset = LensingDataset(val_paths, val_labels)
    test_dataset = LensingDataset(test_paths, test_labels)

    train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
    val_loader = DataLoader(val_dataset, batch_size=batch_size)

```

```

test_loader = DataLoader(test_dataset, batch_size=batch_size)

# Check if dataloaders work
check_data_integrity(train_loader)
except Exception as e:
    print(f"Error creating datasets/dataloaders: {e}")
    return None, 0, 0

# Initialize model
try:
    model = get_model(model_name)
    model = model.to(device)
    print(f"Model initialized: {model_name}")
except Exception as e:
    print(f"Error initializing model: {e}")
    return None, 0, 0

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)

# Train model
try:
    model, history = train_model(
        model, train_loader, val_loader, criterion, optimizer,
        scheduler, num_epochs=num_epochs, device=device
    )
except Exception as e:
    print(f"Error during training: {e}")
    return model, 0, 0

# Evaluate model
try:
    accuracy, mean_auc, plt_roc = evaluate_model(model, test_loader, dev
except Exception as e:
    print(f"Error during evaluation: {e}")
    return model, 0, 0

# Save model to Google Drive
try:
    model_save_path = os.path.join('/content/drive/MyDrive', f'{model_na
    torch.save(model.state_dict(), model_save_path)
    print(f"Model saved to {model_save_path}")
except Exception as e:
    print(f"Error saving model: {e}")

# Plot training history
try:
    plt.figure(figsize=(12, 4))

    plt.subplot(1, 2, 1)
    plt.plot(history['train_loss'], label='Train Loss')
    plt.plot(history['val_loss'], label='Validation Loss')
    plt.title('Loss over Epochs')
    plt.xlabel('Epoch')

```



```

plt.ylabel('Loss')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history['train_acc'], label='Train Accuracy')
plt.plot(history['val_acc'], label='Validation Accuracy')
plt.title('Accuracy over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.savefig('/content/drive/MyDrive/training_history.png')

# Save ROC curve
if plt_roc:
    plt_roc.savefig('/content/drive/MyDrive/roc_curve.png')
except Exception as e:
    print(f"Error plotting results: {e}")

print(f"Final test accuracy: {accuracy:.4f}")
print(f"Final mean AUC: {mean_auc:.4f}")

return model, accuracy, mean_auc

```

```

In [21]: def compare_models(base_dir, models_to_try=['resnet50', 'efficientnet_b0', '
results = {}

for model_name in models_to_try:
    print(f"\n\n{'='*50}")
    print(f"Training model: {model_name}")
    print(f"{'='*50}\n")

    try:
        model, accuracy, mean_auc = main(base_dir, model_name=model_name)
        if model is not None:
            results[model_name] = {'accuracy': accuracy, 'auc': mean_auc}
        else:
            print(f"Training failed for {model_name}")
    except Exception as e:
        print(f"Error in model comparison for {model_name}: {e}")

if not results:
    print("No models were successfully trained")
    return {}

# Print comparison results
print("\nModel Comparison Results:")
print(f"{'Model':<15} {'Accuracy':<10} {'Mean AUC':<10}")
print('-' * 35)

for model_name, metrics in results.items():
    print(f"{'model_name':<15} {metrics['accuracy']:.4f} {metrics['auc']:.4f}")

```

```

# Find best model
best_model = max(results.items(), key=lambda x: x[1]['auc'])
print(f"\nBest model by AUC: {best_model[0]} with AUC = {best_model[1]['auc']}")

# Save results to a text file
try:
    with open('/content/drive/MyDrive/model_comparison_results.txt', 'w') as f:
        f.write("Model Comparison Results:\n")
        f.write(f"{'Model':<15} {'Accuracy':<10} {'Mean AUC':<10}\n")
        f.write('-' * 35 + '\n')

        for model_name, metrics in results.items():
            f.write(f"{'model_name':<15} {'metrics['accuracy']:.4f} {'metrics['mean_auc']:.4f}\n")

        f.write(f"\nBest model by AUC: {best_model[0]} with AUC = {best_model[1]['auc']}")
except Exception as e:
    print(f"Error saving results: {e}")

return results

```

```
In [22]: data_dir = "/content/dataset/dataset/train"
```

```
In [14]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [23]: model, accuracy, mean_auc = main(data_dir, model_name='efficientnet_b0', batch_size=128)
```

```

Using CUDA: Tesla T4
CUDA device count: 1
CUDA version: 12.4
Class distribution:
  Class 0: 10000 samples
  Class 1: 10000 samples
  Class 2: 10000 samples
Train samples: 19200
Validation samples: 4800
Test samples: 6000
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])
Model initialized: efficientnet_b0
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])
Epoch 1/10
-----

```

```
Training: 100%|██████████| 1200/1200 [01:14<00:00, 16.07it/s]
```

Validation: 100%|██████████| 300/300 [00:15<00:00, 19.38it/s]
Val Loss: 0.8092 Acc: 0.6310
Mean AUC: 0.8201
Best validation AUC: 0.8201
Epoch 2/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.52it/s]
Train Loss: 0.0007 Acc: 0.0005
Validation: 100%|██████████| 300/300 [00:05<00:00, 57.54it/s]
Val Loss: 0.6029 Acc: 0.7608
Mean AUC: 0.9168
Best validation AUC: 0.9168
Epoch 3/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.51it/s]
Train Loss: 0.0005 Acc: 0.0007
Validation: 100%|██████████| 300/300 [00:04<00:00, 60.23it/s]
Val Loss: 0.3857 Acc: 0.8435
Mean AUC: 0.9587
Best validation AUC: 0.9587
Epoch 4/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.49it/s]
Train Loss: 0.0006 Acc: 0.0005
Validation: 100%|██████████| 300/300 [00:05<00:00, 59.53it/s]
Val Loss: 0.3235 Acc: 0.8762
Mean AUC: 0.9699
Best validation AUC: 0.9699
Epoch 5/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.54it/s]
Train Loss: 0.0001 Acc: 0.0008
Validation: 100%|██████████| 300/300 [00:04<00:00, 60.99it/s]
Val Loss: 0.3389 Acc: 0.8702
Mean AUC: 0.9697
Best validation AUC: 0.9699
Epoch 6/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.44it/s]
Train Loss: 0.0003 Acc: 0.0007
Validation: 100%|██████████| 300/300 [00:04<00:00, 60.36it/s]
Val Loss: 0.2999 Acc: 0.8931
Mean AUC: 0.9755
Best validation AUC: 0.9755
Epoch 7/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.45it/s]
Train Loss: 0.0001 Acc: 0.0008
Validation: 100%|██████████| 300/300 [00:04<00:00, 60.07it/s]

Val Loss: 0.2799 Acc: 0.9010
Mean AUC: 0.9757
Best validation AUC: 0.9757
Epoch 8/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.66it/s]

Train Loss: 0.0001 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:05<00:00, 55.34it/s]

Val Loss: 0.2332 Acc: 0.9185

Mean AUC: 0.9856

Best validation AUC: 0.9856

Epoch 9/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.65it/s]

Train Loss: 0.0000 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:05<00:00, 53.37it/s]

Val Loss: 0.2263 Acc: 0.9219

Mean AUC: 0.9856

Best validation AUC: 0.9856

Epoch 10/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.65it/s]

Train Loss: 0.0000 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:05<00:00, 52.76it/s]

Val Loss: 0.2383 Acc: 0.9254

Mean AUC: 0.9855

Best validation AUC: 0.9856


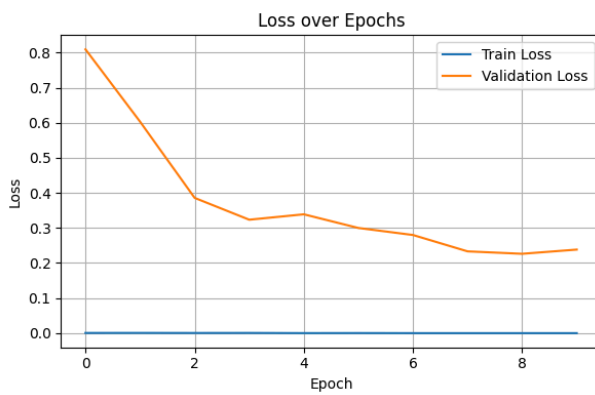

Testing: 100%|██████████| 375/375 [00:15<00:00, 23.93it/s]

Test Accuracy: 0.9223

Model saved to /content/drive/MyDrive/efficientnet_b0_lensing_classifier.pth

Final test accuracy: 0.9223

Final mean AUC: 0.9856



```
In [24]: results = compare_models(data_dir, models_to_try=['resnet50', 'efficientnet_
```

```
=====
Training model: resnet50
=====
```

Using CUDA: Tesla T4

CUDA device count: 1

CUDA version: 12.4

Class distribution:

Class 0: 10000 samples

Class 1: 10000 samples

Class 2: 10000 samples

Train samples: 19200

Validation samples: 4800

Test samples: 6000

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth

100%|██████████| 97.8M/97.8M [00:00<00:00, 172MB/s]

Model initialized: resnet50

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Epoch 1/10

Training: 100%|██████████| 1200/1200 [01:52<00:00, 10.70it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.76it/s]

Val Loss: 1.1029 Acc: 0.3310

Mean AUC: 0.5027

Best validation AUC: 0.5027

Epoch 2/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.77it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.33it/s]

Val Loss: 1.0991 Acc: 0.3333

Mean AUC: 0.5045

Best validation AUC: 0.5045

Epoch 3/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.75it/s]

Train Loss: 0.0009 Acc: 0.0002

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.34it/s]

Val Loss: 1.0987 Acc: 0.3356
Mean AUC: 0.5047
Best validation AUC: 0.5047
Epoch 4/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.77it/s]

Train Loss: 0.0009 Acc: 0.0002

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.71it/s]

Val Loss: 1.1028 Acc: 0.3233

Mean AUC: 0.4941

Best validation AUC: 0.5047

Epoch 5/10

Training: 100%|██████████| 1200/1200 [01:50<00:00, 10.82it/s]

Train Loss: 0.0009 Acc: 0.0001

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.71it/s]

Val Loss: 1.1056 Acc: 0.3381

Mean AUC: 0.4925

Best validation AUC: 0.5047

Epoch 6/10

Training: 100%|██████████| 1200/1200 [01:50<00:00, 10.82it/s]

Train Loss: 0.0009 Acc: 0.0005

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.56it/s]

Val Loss: 1.0983 Acc: 0.3471

Mean AUC: 0.5148

Best validation AUC: 0.5148

Epoch 7/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.79it/s]

Train Loss: 0.0009 Acc: 0.0002

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.84it/s]

Val Loss: 1.0988 Acc: 0.3281

Mean AUC: 0.4996

Best validation AUC: 0.5148

Epoch 8/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.78it/s]

Train Loss: 0.0009 Acc: 0.0003

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.63it/s]

Val Loss: 1.0997 Acc: 0.3331

Mean AUC: 0.4945

Best validation AUC: 0.5148

Epoch 9/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.80it/s]

Train Loss: 0.0009 Acc: 0.0003

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.88it/s]

Val Loss: 1.0990 Acc: 0.3346
Mean AUC: 0.4865
Best validation AUC: 0.5148
Epoch 10/10

Training: 100%|██████████| 1200/1200 [01:51<00:00, 10.79it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.87it/s]

Val Loss: 1.0988 Acc: 0.3281

Mean AUC: 0.4943

Best validation AUC: 0.5148

Testing: 100%|██████████| 375/375 [00:11<00:00, 31.91it/s]

Test Accuracy: 0.3407

Model saved to /content/drive/MyDrive/resnet50_lensing_classifier.pth

Final test accuracy: 0.3407

Final mean AUC: 0.5092

=====
Training model: efficientnet_b0
=====

Using CUDA: Tesla T4

CUDA device count: 1

CUDA version: 12.4

Class distribution:

Class 0: 10000 samples

Class 1: 10000 samples

Class 2: 10000 samples

Train samples: 19200

Validation samples: 4800

Test samples: 6000

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Model initialized: efficientnet_b0

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])

Data check - Input dtype: torch.float32, Labels dtype: torch.int64

Data check - Labels unique values: tensor([0, 1, 2])

Epoch 1/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.22it/s]

Train Loss: 0.0008 Acc: 0.0005

Validation: 100%|██████████| 300/300 [00:05<00:00, 53.72it/s]

Val Loss: 0.8767 Acc: 0.5575

Mean AUC: 0.8040

Best validation AUC: 0.8040

Epoch 2/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.34it/s]

Train Loss: 0.0004 Acc: 0.0007

Validation: 100%|██████████| 300/300 [00:05<00:00, 56.46it/s]

Val Loss: 0.5076 Acc: 0.7950

Mean AUC: 0.9326

Best validation AUC: 0.9326

Epoch 3/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.17it/s]

Train Loss: 0.0005 Acc: 0.0006

Validation: 100%|██████████| 300/300 [00:05<00:00, 56.48it/s]

Val Loss: 0.3957 Acc: 0.8494

Mean AUC: 0.9558

Best validation AUC: 0.9558

Epoch 4/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.22it/s]

Train Loss: 0.0005 Acc: 0.0006

Validation: 100%|██████████| 300/300 [00:05<00:00, 58.30it/s]

Val Loss: 0.3101 Acc: 0.8783

Mean AUC: 0.9717

Best validation AUC: 0.9717

Epoch 5/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.40it/s]

Train Loss: 0.0001 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:04<00:00, 61.00it/s]

Val Loss: 0.3444 Acc: 0.8696

Mean AUC: 0.9738

Best validation AUC: 0.9738

Epoch 6/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.42it/s]

Train Loss: 0.0004 Acc: 0.0007

Validation: 100%|██████████| 300/300 [00:04<00:00, 60.51it/s]

Val Loss: 0.2732 Acc: 0.8969

Mean AUC: 0.9779

Best validation AUC: 0.9779

Epoch 7/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.22it/s]

Train Loss: 0.0003 Acc: 0.0007

Validation: 100%|██████████| 300/300 [00:05<00:00, 59.32it/s]

Val Loss: 0.2608 Acc: 0.9012

Mean AUC: 0.9795

Best validation AUC: 0.9795

Epoch 8/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.25it/s]

Train Loss: 0.0000 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:05<00:00, 59.57it/s]

Val Loss: 0.2266 Acc: 0.9198
Mean AUC: 0.9859
Best validation AUC: 0.9859
Epoch 9/10

Training: 100%|██████████| 1200/1200 [00:59<00:00, 20.34it/s]
Train Loss: 0.0001 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:04<00:00, 62.23it/s]
Val Loss: 0.2285 Acc: 0.9206
Mean AUC: 0.9861
Best validation AUC: 0.9861
Epoch 10/10

Training: 100%|██████████| 1200/1200 [00:58<00:00, 20.63it/s]
Train Loss: 0.0000 Acc: 0.0008

Validation: 100%|██████████| 300/300 [00:05<00:00, 56.92it/s]
Val Loss: 0.2378 Acc: 0.9260
Mean AUC: 0.9862
Best validation AUC: 0.9862

Testing: 100%|██████████| 375/375 [00:06<00:00, 61.41it/s]
Test Accuracy: 0.9267
Model saved to /content/drive/MyDrive/efficientnet_b0_lensing_classifier.pth
Final test accuracy: 0.9267
Final mean AUC: 0.9859

=====
Training model: densenet121
=====

Using CUDA: Tesla T4
CUDA device count: 1
CUDA version: 12.4
Class distribution:
 Class 0: 10000 samples
 Class 1: 10000 samples
 Class 2: 10000 samples
Train samples: 19200
Validation samples: 4800
Test samples: 6000
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])

Downloading: "https://download.pytorch.org/models/densenet121-a639ec97.pth"
to /root/.cache/torch/hub/checkpoints/densenet121-a639ec97.pth
100%|██████████| 30.8M/30.8M [00:00<00:00, 82.4MB/s]

Model initialized: densenet121
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])
Data check - Input shape: torch.Size([16, 3, 150, 150]), Labels shape: torch.Size([16])
Data check - Input dtype: torch.float32, Labels dtype: torch.int64
Data check - Labels unique values: tensor([0, 1, 2])
Epoch 1/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.20it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 32.55it/s]

Val Loss: 1.1730 Acc: 0.3333

Mean AUC: 0.5079

Best validation AUC: 0.5079

Epoch 2/10

Training: 100%|██████████| 1200/1200 [01:46<00:00, 11.29it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 32.30it/s]

Val Loss: 1.1325 Acc: 0.3317

Mean AUC: 0.5003

Best validation AUC: 0.5079

Epoch 3/10

Training: 100%|██████████| 1200/1200 [01:46<00:00, 11.30it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.25it/s]

Val Loss: 1.1644 Acc: 0.3333

Mean AUC: 0.5114

Best validation AUC: 0.5114

Epoch 4/10

Training: 100%|██████████| 1200/1200 [01:46<00:00, 11.26it/s]

Train Loss: 0.0009 Acc: 0.0002

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.10it/s]

Val Loss: 1.1629 Acc: 0.3321

Mean AUC: 0.4953

Best validation AUC: 0.5114

Epoch 5/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.20it/s]

Train Loss: 0.0009 Acc: 0.0003

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.32it/s]

Val Loss: 1.1045 Acc: 0.3333

Mean AUC: 0.4956

Best validation AUC: 0.5114

Epoch 6/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.14it/s]

Loading [MathJax]/extensions/Safe.js } 0.0009 Acc: 0.0002

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.25it/s]

Val Loss: 6.2354 Acc: 0.3333

Mean AUC: 0.5017

Best validation AUC: 0.5114

Epoch 7/10

Training: 100%|██████████| 1200/1200 [01:46<00:00, 11.22it/s]

Train Loss: 0.0009 Acc: 0.0003

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.24it/s]

Val Loss: 1.1089 Acc: 0.3333

Mean AUC: 0.5062

Best validation AUC: 0.5114

Epoch 8/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.21it/s]

Train Loss: 0.0009 Acc: 0.0001

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.83it/s]

Val Loss: 1.1016 Acc: 0.3358

Mean AUC: 0.5019

Best validation AUC: 0.5114

Epoch 9/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.17it/s]

Train Loss: 0.0009 Acc: 0.0003

Validation: 100%|██████████| 300/300 [00:09<00:00, 31.28it/s]

Val Loss: 1.0985 Acc: 0.3383

Mean AUC: 0.5073

Best validation AUC: 0.5114

Epoch 10/10

Training: 100%|██████████| 1200/1200 [01:47<00:00, 11.12it/s]

Train Loss: 0.0009 Acc: 0.0004

Validation: 100%|██████████| 300/300 [00:09<00:00, 30.72it/s]

Val Loss: 1.0991 Acc: 0.3477

Mean AUC: 0.5026

Best validation AUC: 0.5114

Testing: 100%|██████████| 375/375 [00:12<00:00, 31.12it/s]

Test Accuracy: 0.3388

Model saved to /content/drive/MyDrive/densenet121_lensing_classifier.pth



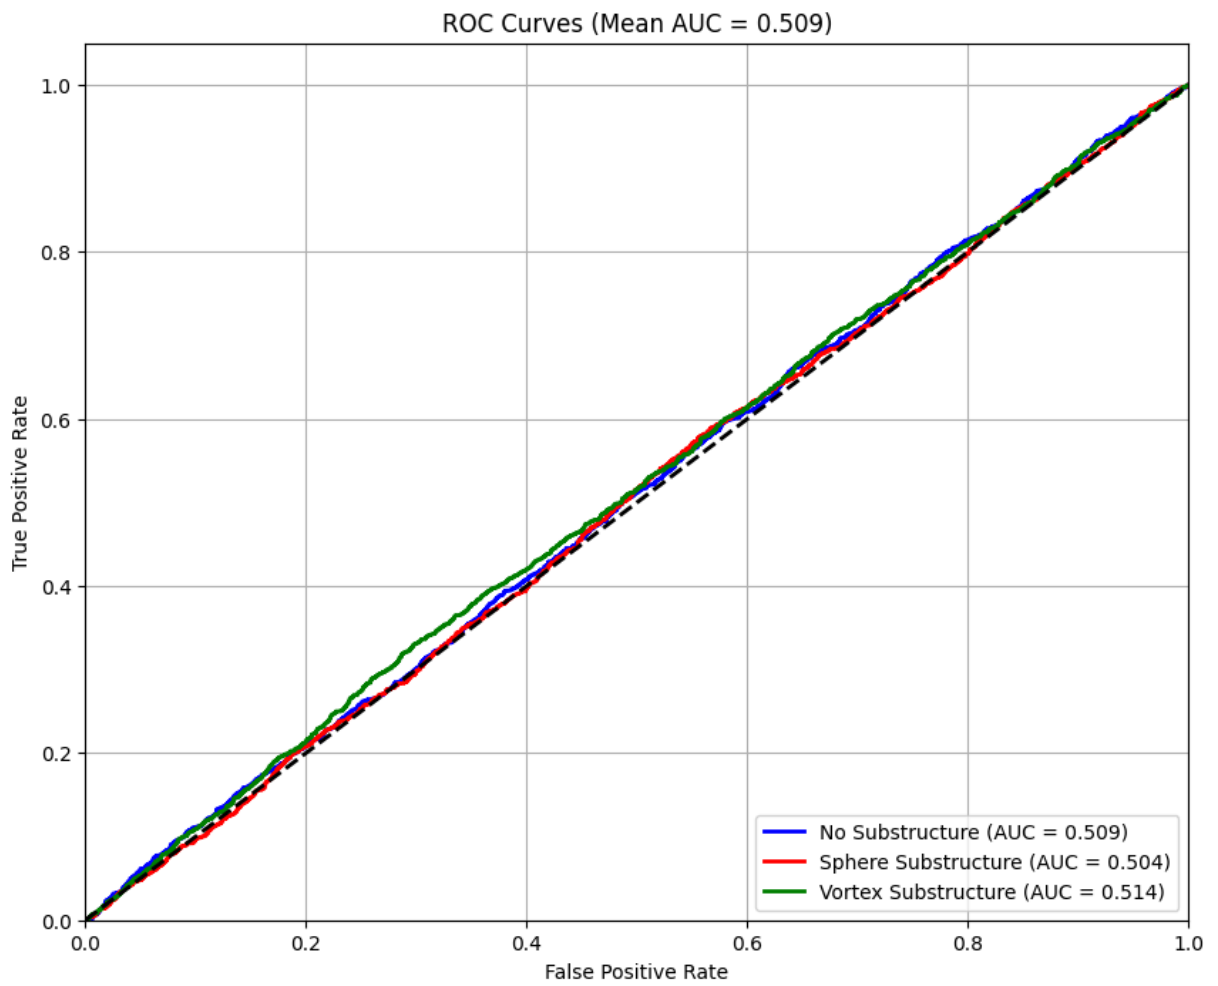
Final test accuracy: 0.3388




Final mean AUC: 0.5093

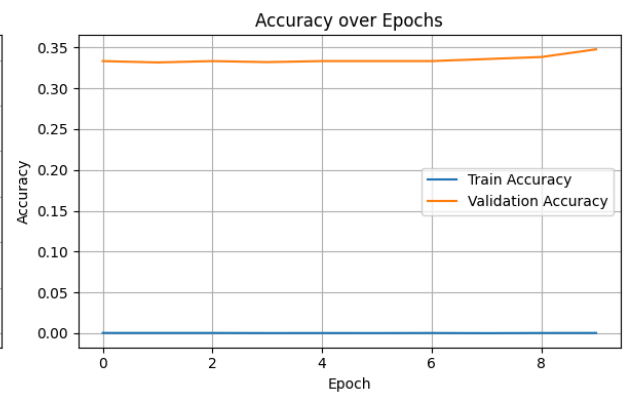

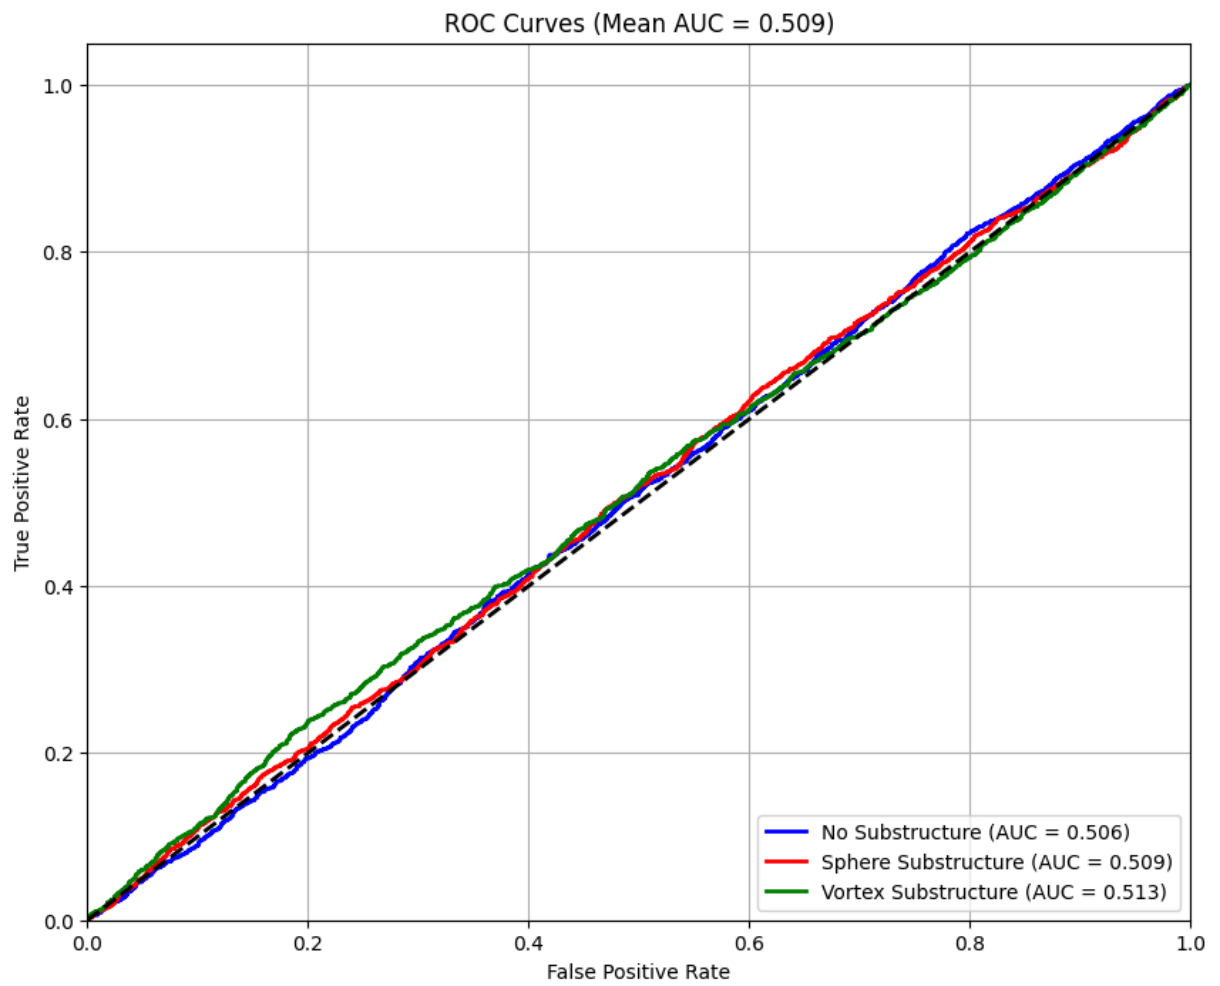
Model Comparison Results:

Model	Accuracy	Mean AUC
resnet50	0.3407	0.5092
efficientnet_b0	0.9267	0.9859
densenet121	0.3388	0.5093

Best model by AUC: efficientnet_b0 with AUC = 0.9859







This notebook was converted with convert.ploomber.io